Classifying Melanoma

What Is the Melanoma Dataset?

- Annual competition on Kaggle
  - Lots of highly skilled veterans!
- Expert Judgement: 74.1% sensitivity and 60% specificity by dermatologists
- Challenges
  - Large Data Set: 35 GB+ of JPEGs
  - Unbalanced classes: Only 500 positive examples out of n=30,000+)
  - Heterogeneous format: Inconsistent image quality/shape/brightness/orientation
  - Visually ambiguous classes

Can you tell which one is Melanoma?

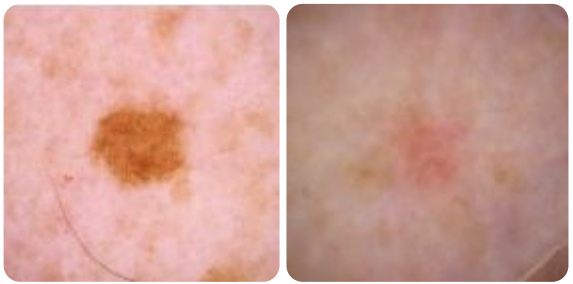

Sexy hair...

Round crop!?

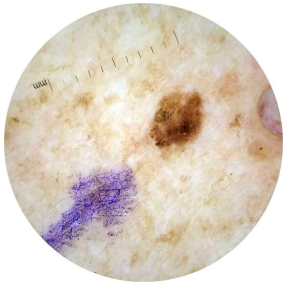

Baseline Models

Baseline Models

- Logistic Regression
- Support Vector Machine

Baseline Results

Logistic Regression

	precision	recall	f1-score	support
0	0.98	1.00	0.99	6441
1	0.00	0.00	0.00	117
accuracy			0.98	6558
macro avg	0.49	0.50	0.50	6558
weighted avg	0.96	0.98	0.97	6558

SVM (Linear Kernel)

	precision	recall	f1-score	support
0	0.98	0.99	0.99	6440
1	0.07	0.05	0.06	116
accuracy			0.97	6556
macro avg	0.53	0.52	0.52	6556
weighted avg	0.97	0.97	0.97	6556

Notes:

- 10 Epochs
- Sigmoid Threshold = 0.5
- No Regularization

Neural Networks

Three Attempts

- V1: CNN Barebone
- V2: CNN Advanced
- V3: CNN (Image) + Neural Network (Metadata)

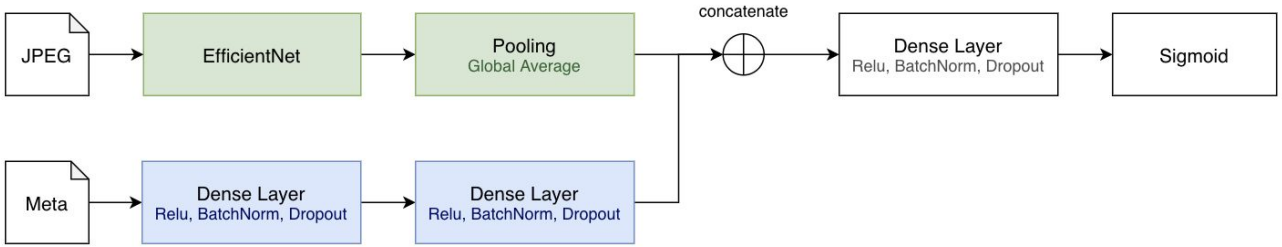

- Bonus EDA: Can neural network learn?

V1: CNN Barebone

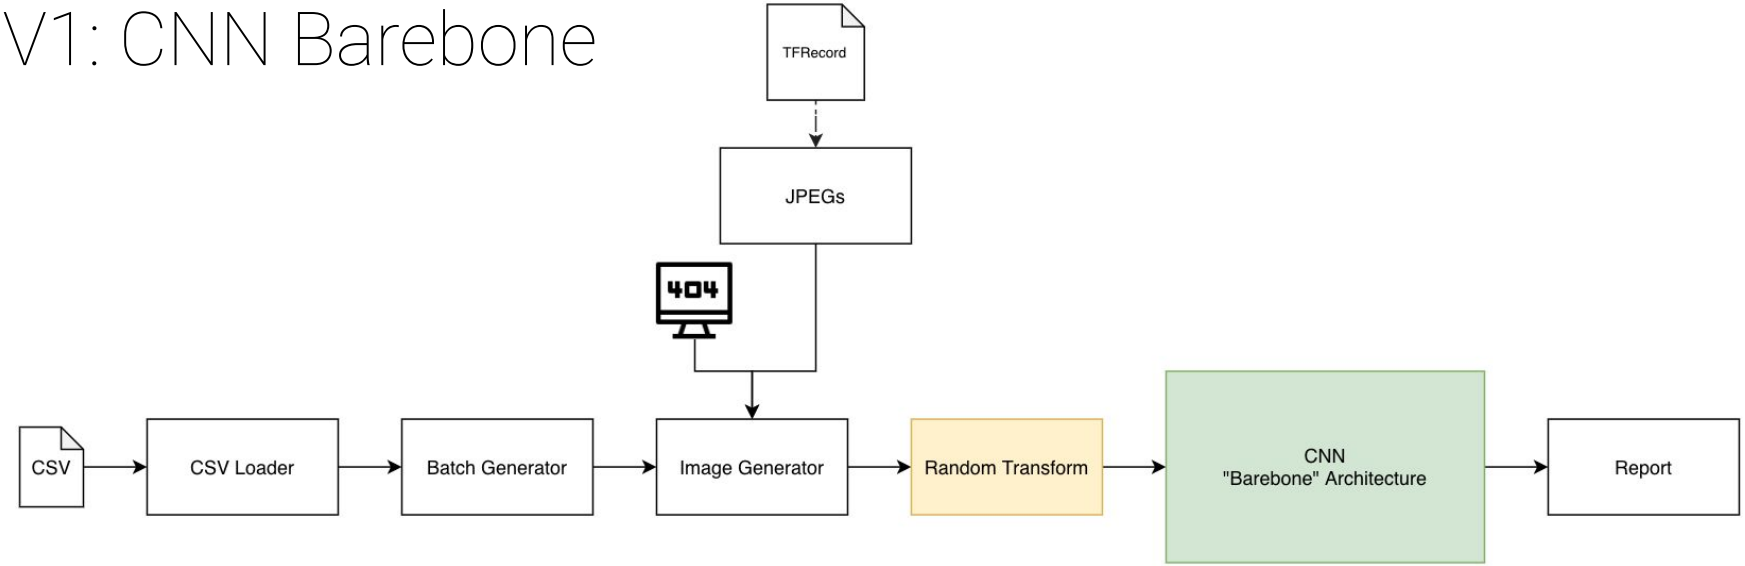

Architecture 1

- Input
- Conv. 3x3x64, ReLu
- Conv. 3x3x32, ReLu
- Flatten
- Dense, Softmax

Architecture 2

- Input
- Conv. 3x3x75, ReLu
  - Pool 2x2
  - Dropout 0.15
- Conv. 3x3x125, ReLu
  - Pool 2x2
  - Dropout 0.15
- Flatten
- Dense, Softmax

V1: Miserable Failure

First run...

	precision	recall	f1-score
0	0.50	1.00	0.67
1	0.00	0.00	0.00
accuracy			0.50
macro avg	0.25	0.50	0.33
weighted avg	0.25	0.50	0.33

After countless hours of tweaking...

	precision	recall	f1-score
0	0.53	0.96	0.68
1	0.78	0.14	0.24
accuracy			0.55
macro avg	0.65	0.55	0.46
weighted avg	0.65	0.55	0.46

And... a full epoch takes ~1 hour to complete...
Scrapped 1,000+ lines of codes... Ouch!

V2: CNN Advanced

TPU + TFRecord + K-Fold = Efficient Training

Stratified Training Data	EfficientNet	Transfer Learning	Focal Loss

Others

- Additional feature engineering
- Apply random transformation on both training and test datasets
- Restrict image augmentation (e.g. avoid black edges)
- Experiment with hyperparameters using k-fold CV

Stratified Training Data & K-Fold

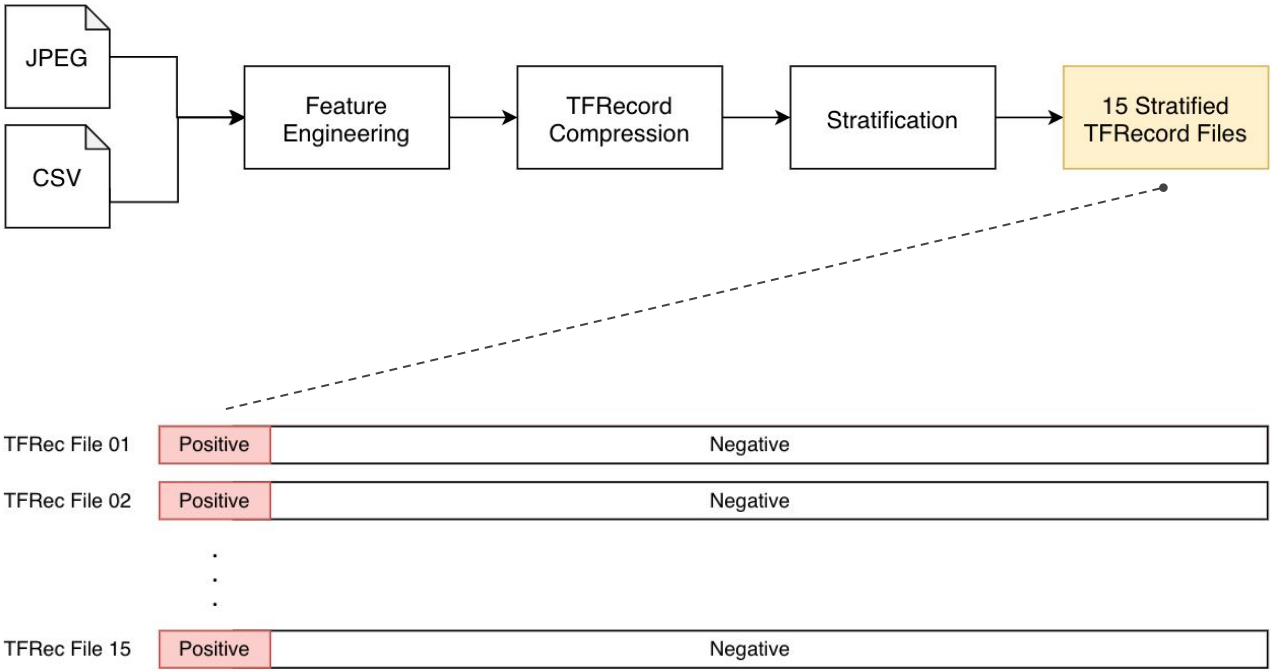

What Is TPU?

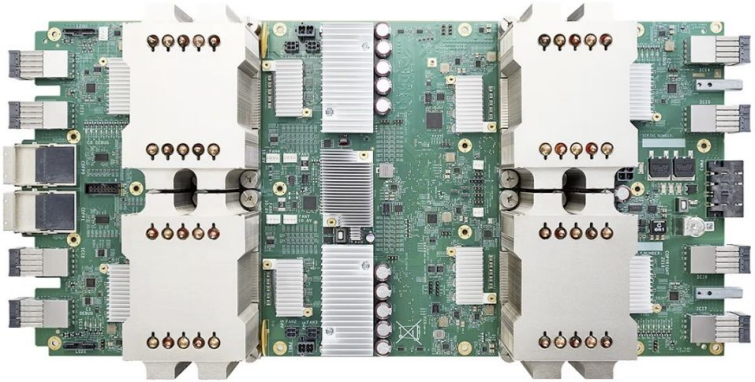

colab

- 4 chips, each has 2 cores
- Processor designed for NN & Tensorflow
- GCS & code changes required (e.g. batch size)!

Colab & TPU

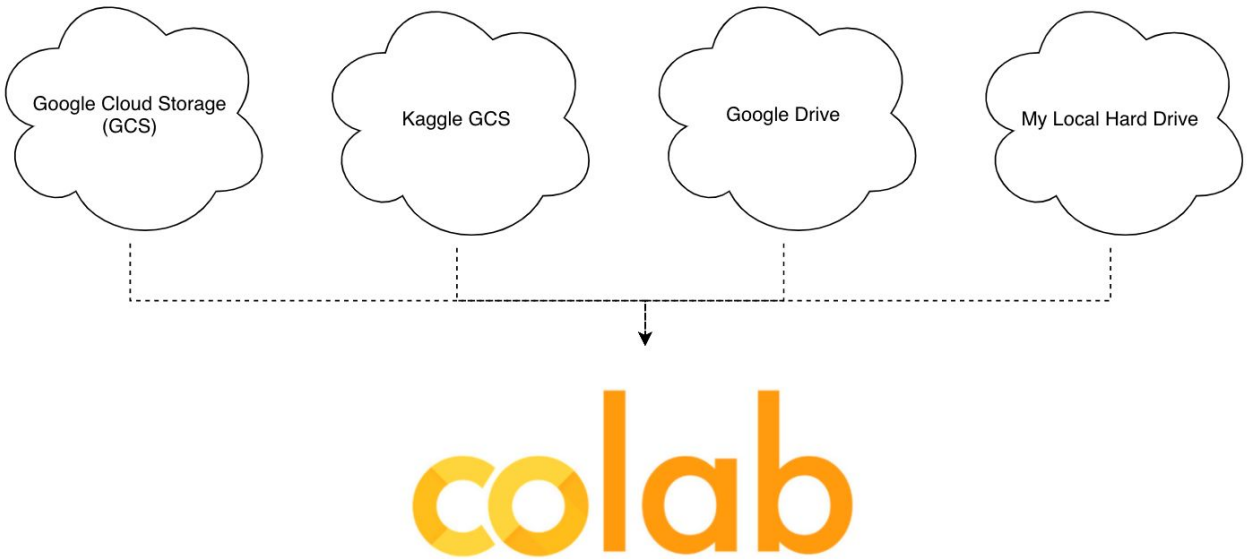
Google Cloud Storage (GCS)
Kaggle GCS
Google Drive
My Local Hard Drive
colab

EfficientNet

Problem: How do we scale up a CNN?

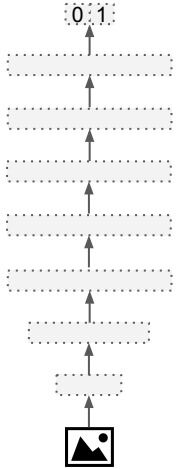

Deeper

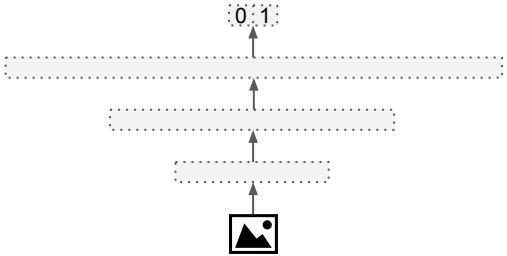

Wider

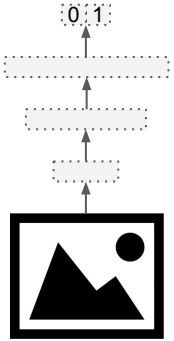

Higher Resolution

EfficientNet

- Baseline Network
  - EfficientNet-B0 is constructed using neural architecture search, which optimizes both accuracy and efficiency (FLOPS).
- Compound Scaling
  - Grid search to optimally scale up EfficientNet-B0 under a fixed resource constraint
  - Result: B0, B1, B2, ..., B7

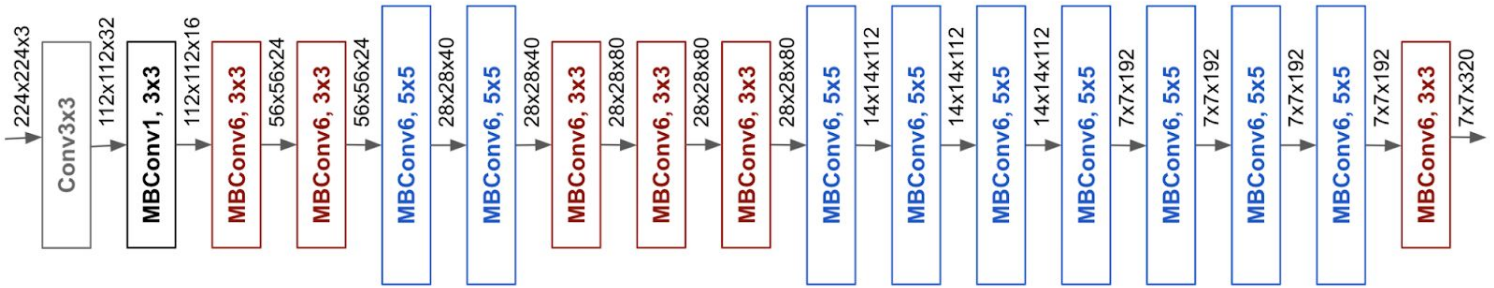

EfficientNet-B0

Transferred Learning

EfficientNet
(Trained by Google)

Transfer Weights

EfficientNet
(My Model)

PRE-TRAINED

IMAGENET

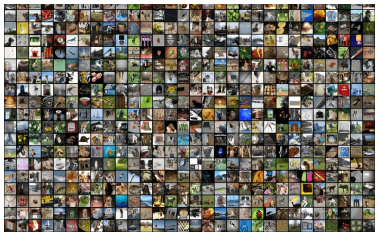

TRAINING CAT PICTURES

kaggle

Focal Loss: ~~Finding~~ Learning Waldo

What if... for every positive example, you have 100+ negative examples?

Solution: Weight Adjustment

1 Heavy-Weight Waldo

False Negative

100+ Small Strangers

Correct Predictions

What Is Focal Loss?

$$FL = -\sum_{i=1}^{C=2} (1 - s_i)^{\gamma} t_i \log(s_i)$$

C is the class being predicted
Si is the probability predicted by the neural network

Hyperparameters

- t_i (α on FB's paper) is the probability for positive case
  - 1 in the ideal world
  - Can be adjusted to offset the imbalance in the dataset
- γ is used to adjust the loss curve
  - 2 is the proposed default by FB

Binary Cross-Entropy

$$CE = -\sum_{i=1}^{C=2} t_i \log(s_i)$$

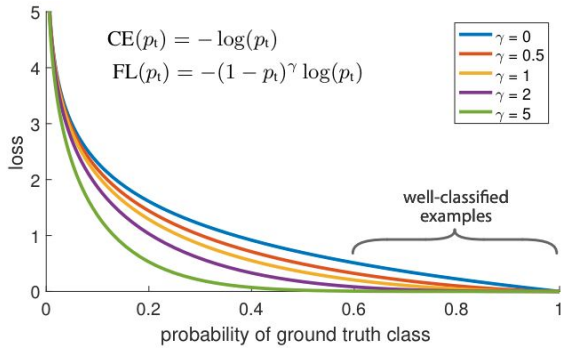

$$L = -\alpha_t (1 - p_t)^{\gamma} \log(p_t)$$

$$\alpha_t = \begin{cases} \alpha & \text{if } y = 1 \\ 1 - \alpha & \text{otherwise} \end{cases}$$

V2 Architecture

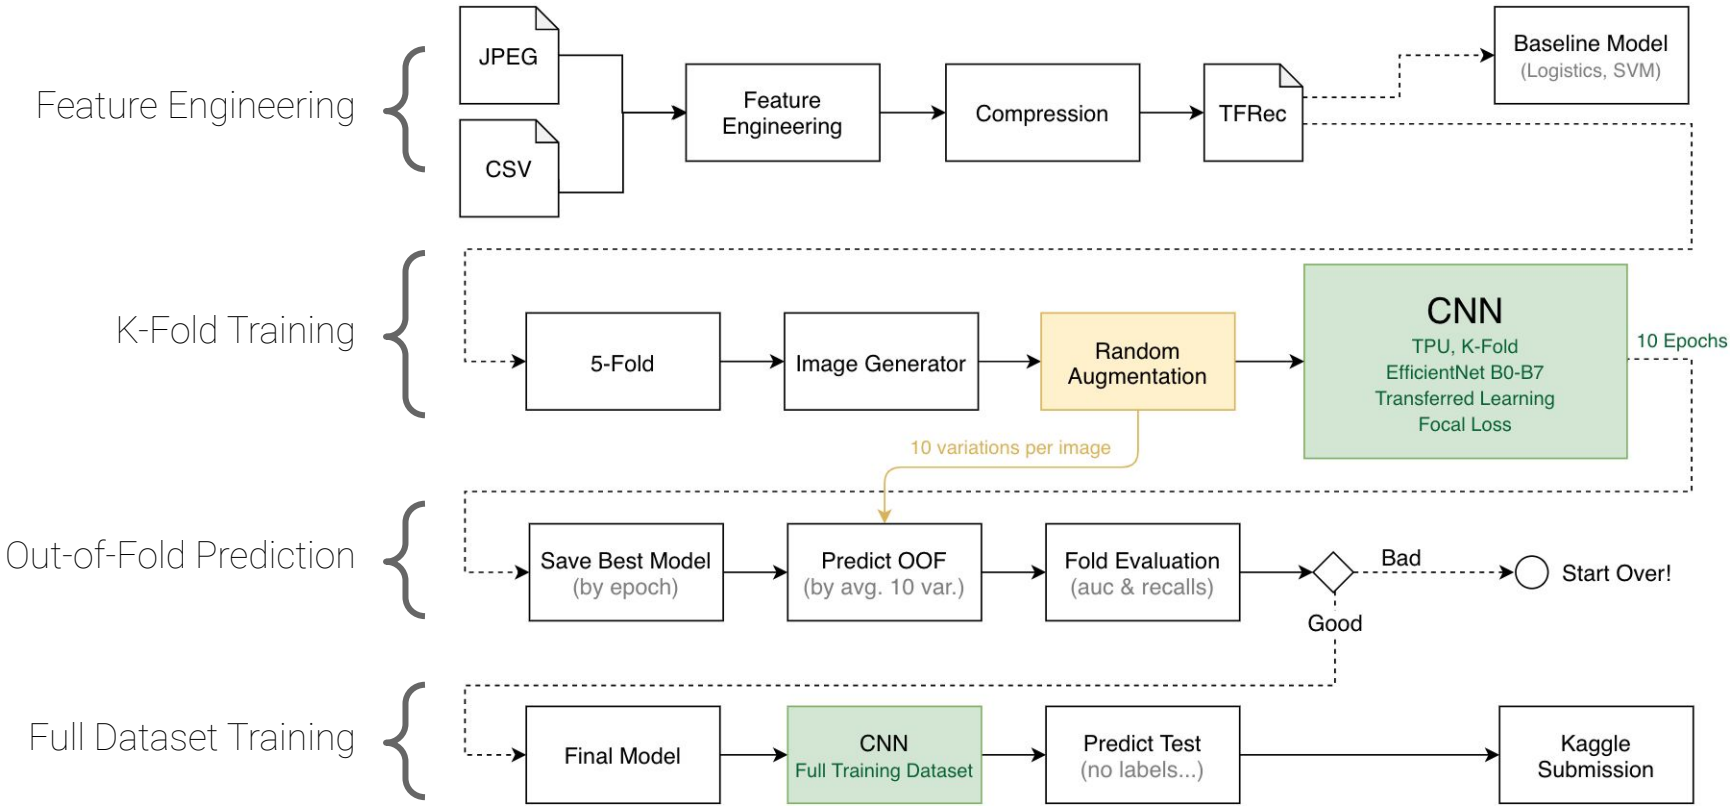

V3 Architecture

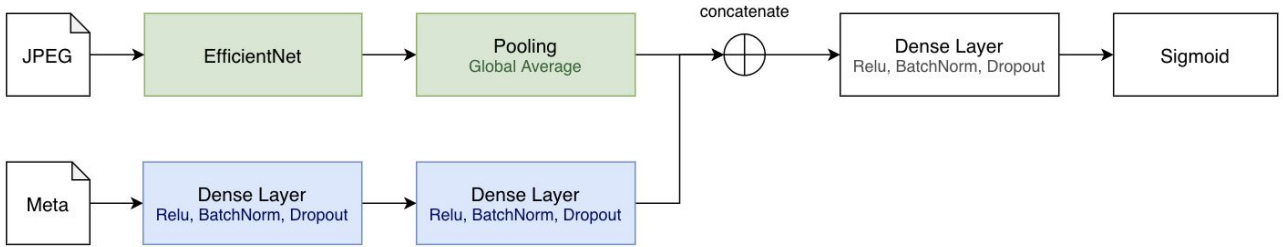

V3 Predictions

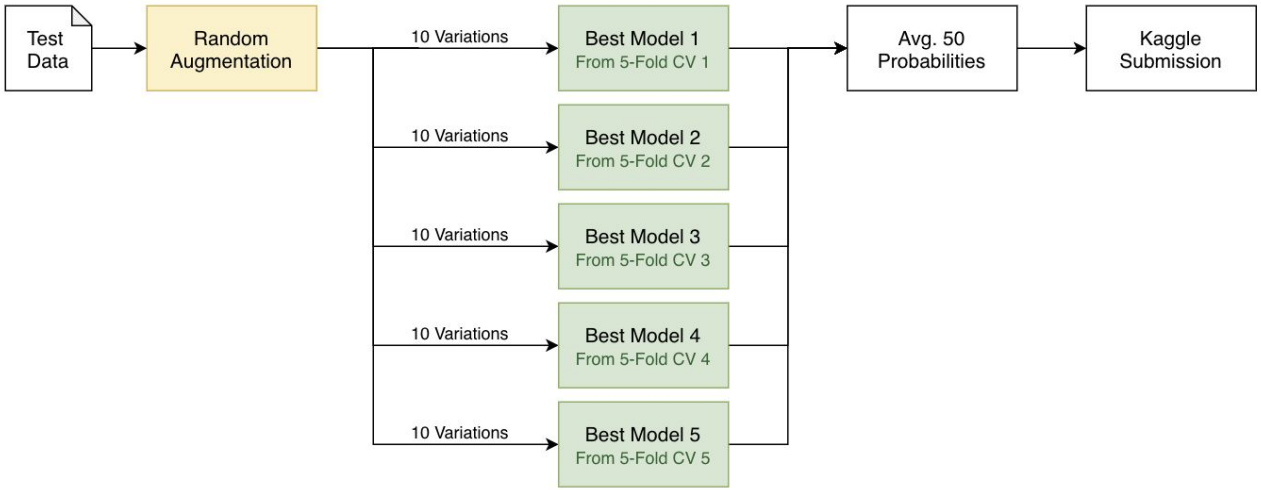

Results

V2 Results

Rank #1, 2019 Competition

Configuration	Sensitivity	AUC
SENet154 SS 224×224	67.2 ± 0.8	94.1 ± 0.4
EN B0 SS 224×224 $C = 8$	66.7 ± 1.8	94.0 ± 0.3
EN B0 SS 224×224	65.8 ± 1.7	92.9 ± 0.8
EN B0 RR 224×224	67.0 ± 1.6	93.3 ± 0.5
EN B1 SS 240×240	65.9 ± 1.6	93.8 ± 0.3
EN B1 RR 240×240	66.8 ± 1.5	93.7 ± 0.7
EN B2 SS 260×260	67.2 ± 1.4	94.1 ± 0.6
EN B2 RR 260×260	67.6 ± 2.0	93.9 ± 0.7
EN B3 SS 300×300	67.8 ± 2.0	94.3 ± 0.5
EN B3 RR 300×300	67.0 ± 1.5	94.0 ± 0.6
EN B4 SS 380×380	67.8 ± 1.1	94.1 ± 0.5
EN B4 RR 380×380	68.1 ± 1.6	94.2 ± 0.5
EN B5 SS 456×456	68.2 ± 0.9	93.5 ± 0.7
EN B5 RR 456×456	68.0 ± 2.2	94.0 ± 0.7
EN B6 SS 456×456	68.8 ± 0.7	93.8 ± 0.5
Ensemble Average	71.7 ± 1.7	95.4 ± 0.5
Ensemble Optimal	$\mathbf{72.5 \pm 1.7}$	$\mathbf{95.4 \pm 0.5}$

Our Results

AUC: 89.05

Problem: Underfitting (bias)

- Training and validation suffers from the same performance issue
- ~0.8 AUC, ~0.2 Recall

Rank: 1,498 (out of 2,600+ entries)

Configs:

- Model: EfficientNet B7
- Loss Function: Focal Loss
- Transferred Learning: ImageNet
- Epochs: 8
- Batch Size: 64
- Image Size: 256x256
- Random Image Augmentation: True

V3 Results

Problem: Overfitting (high variance), large performance gap between training and validation

Epoch 5:

	Training Data	Validation Data
AUC	1.0	0.71
Recall	100%	0%

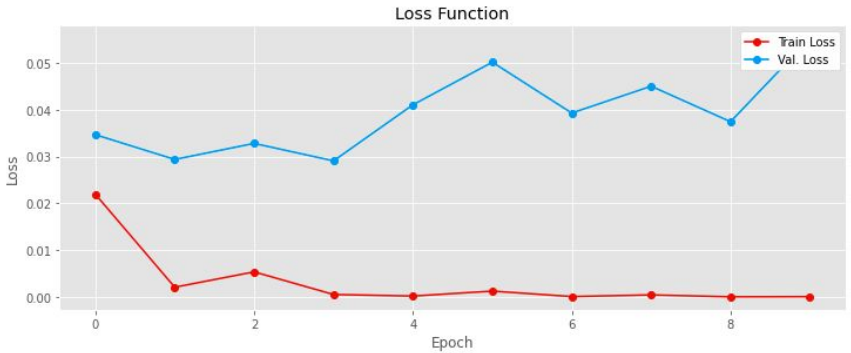

Demo

What Else Can Be Done?

v2 (underfitting) and v3 (overfitting) suffer from different issues and require different treatments.

Other Things to Try:

- Ensemble!
- Hyperparameters Tuning
  - Gradient Clipping
  - Learning Rate Decay
  - Grid Search
- Feature Engineering
  - Hair Simulation
  - Microscope Vignette Simulation
  - Higher Resolutions

Appendix

V3 TensorFlow Model

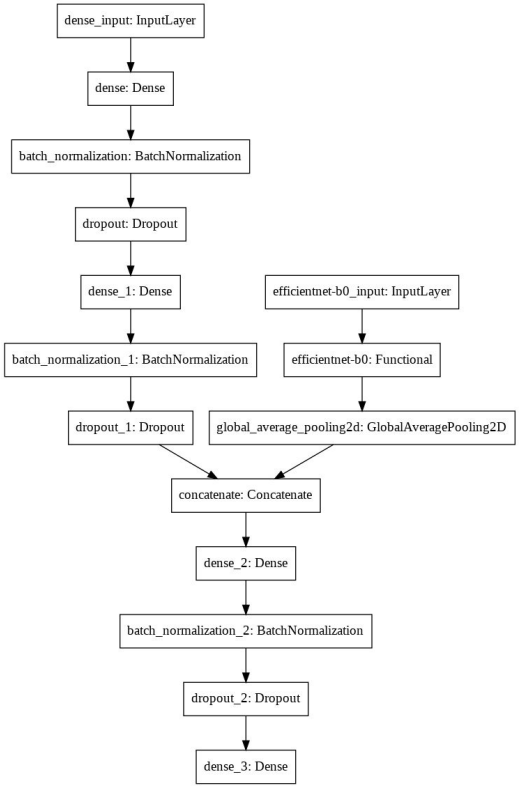

References

EfficientNet: Improving Accuracy and Efficiency through AutoML and Model Scaling
https://tinyurl.com/y6a78p98

What makes TPUs fine-tuned for deep learning?
https://tinyurl.com/yxqw5x4e

Understanding Categorical Cross-Entropy Loss, Binary Cross-Entropy Loss, Softmax Loss, Logistic Loss, Focal Loss and all those confusing names
https://tinyurl.com/y6x5n9w5

Focal Loss for Dense Object Detection
https://tinyurl.com/y2xua29r